


```
# Lab04 - Statistical Data Analytics
## Nguyen Quoc Tuan - 19522476
## Link github: https://github.com/tuNqws/data\_mining.git
```

```
from sklearn import datasets
import matplotlib.pyplot as plt
```

```
digits = datasets.load_digits()
```

```
plt.imshow(digits.images[1010], cmap=plt.cm.gray_r, interpolation='nearest')
plt.show()
```



```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from urllib.request import urlopen
```

```
df = pd.read_csv('gapminder.csv')
sns.heatmap(df.corr(), square=True, cmap='RdYlGn')
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.show()
```

```
<ipython-input-13-6be5d627e500>:2: FutureWarning: The default value of numeric_only in c
sns.heatmap(df.corr(), square=True, cmap='RdYlGn')
```



```
# modified/added by Jinny
plt.style.use('ggplot')

y = df['life'].values
X = df.drop('life', axis=1)

# Reshape to 1-D
y = y.reshape(-1, 1)
X_fertility = X['fertility'].values.reshape(-1, 1)

_ = plt.scatter(X['fertility'], y, color='blue')
_ = plt.ylabel('Life Expectancy')
_ = plt.xlabel('Fertility')

# -----
# Import LinearRegression
from sklearn.linear_model import LinearRegression

# Create the regressor: reg
reg = LinearRegression()

# Create the prediction space
prediction_space = np.linspace(min(X_fertility), max(X_fertility)).reshape(-1,1)


# Fit the model to the data
reg.fit(X_fertility, y)

# Compute predictions over the prediction space: y_pred
y_pred = reg.predict(prediction_space)

# Print R^2
print(reg.score(X_fertility, y))

# Plot regression line
plt.plot(prediction_space, y_pred, color='black', linewidth=3)
plt.show()
```

0.6192442167740035



```
from sklearn.model_selection import train_test_split
features = pd.read_csv('gapminder.csv')
df = pd.read_csv('gapminder.csv')
del features['life']
del features['Region']

y_life = df['life'].values.reshape(-1,1)

# Create training and test sets
X_train, X_test, y_train, y_test = train_test_split(features, y_life, test_size = 0.3, random_state=42)
```

```
# Create the regressor: reg_all
reg_all = LinearRegression()

# Fit the regressor to the training data
reg_all.fit(X_train, y_train)


# Predict on the test data: y_pred
```

```
# Compute and print R^2 and RMSE
print(reg_all.score(features, y_life))
```

```
0.8914651485793176
```

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
N =10
m = .9
c = 1
x = np.linspace(0,2*np.pi,N)
y=m*x + c+ np.random.normal(0,.3,x.shape)
plt.figure()
plt.plot(x,y,'o')
plt.xlabel('x')
plt.ylabel('y')
plt.show()
```



```
import torch
```

```
from torch.utils.data import Dataset
class MyDataset(Dataset):
    def __init__(self, x, y):
        self.x=x
        self.y=y

    def __len__(self):
        return len(self.x)


    def __getitem__(self, idx):
        sample = {
            'feature': torch.tensor ([1, self.x[idx]]),
            'label': torch.tensor ([self.y[idx]])}
        return sample
```

```
dataset = MyDataset(x, y)
for i in range(len(dataset)):
    sample = dataset[i]
    print(i, sample['feature'], sample['label'])

0 tensor([1., 0.], dtype=torch.float64) tensor([1.3249], dtype=torch.float64)
1 tensor([1.0000, 0.6981], dtype=torch.float64) tensor([1.6252], dtype=torch.float64)
2 tensor([1.0000, 1.3963], dtype=torch.float64) tensor([2.5060], dtype=torch.float64)
3 tensor([1.0000, 2.0944], dtype=torch.float64) tensor([3.1148], dtype=torch.float64)
4 tensor([1.0000, 2.7925], dtype=torch.float64) tensor([3.4458], dtype=torch.float64)
5 tensor([1.0000, 3.4907], dtype=torch.float64) tensor([3.9926], dtype=torch.float64)
6 tensor([1.0000, 4.1888], dtype=torch.float64) tensor([4.6896], dtype=torch.float64)
7 tensor([1.0000, 4.8869], dtype=torch.float64) tensor([5.3221], dtype=torch.float64)
8 tensor([1.0000, 5.5851], dtype=torch.float64) tensor([6.2227], dtype=torch.float64)
9 tensor([1.0000, 6.2832], dtype=torch.float64) tensor([7.3215], dtype=torch.float64)
```

```
from torch.utils.data import DataLoader
dataset= MyDataset(x, y)
batch_size = 4
shuffle = True
num_workers = 4
dataloader = DataLoader( dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_workers )
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:561: UserWarning: This DataLoader will create 4 worker processes
warnings.warn(_create_warning_msg(
```



```
import pprint as pp
for i_batch, samples in enumerate (dataloader):
    print('\nbatch# = %s' % i_batch)
    print('samples: ')
    pp.pprint(samples)

batch# = 0
samples:
{'feature': tensor([[1.0000, 1.3963],
                    [1.0000, 2.7925],
                    [1.0000, 4.1888],
                    [1.0000, 6.2832]], dtype=torch.float64),
 'label': tensor([[2.5060],
                  [3.4458],
                  [4.6896],
                  [7.3215]], dtype=torch.float64)}

batch# = 1
samples:
{'feature': tensor([[1.0000, 2.0944],
                    [1.0000, 3.4907],
                    [1.0000, 0.6981],
                    [1.0000, 4.8869]], dtype=torch.float64),
 'label': tensor([[3.1148],
                  [3.9926],
                  [1.6252],
                  [5.3221]], dtype=torch.float64)}

batch# = 2
samples:
{'feature': tensor([[1.0000, 5.5851],
                    [1.0000, 0.0000]], dtype=torch.float64),
 'label': tensor([[6.2227],
                  [1.3249]], dtype=torch.float64)}
```

```
import torch.nn as nn
import torch.nn.functional as F
class MyModel(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(MyModel, self).__init__()
        self.linear = nn.Linear(input_dim, output_dim)
    def forward (self, x):
        out = self.linear(x)
        return out
```

```
input_dim = 2
output_dim = 1
```

```
model = MyModel(input_dim, output_dim)
```

```

cost = nn.MSELoss()

num_epochs = 10 # How many times the entire training data is seen?
l_rate = 0.01
optimiser = torch.optim.SGD(model.parameters(), lr = l_rate)

dataset = MyDataset(torch.Tensor(x), torch.Tensor(y)) # Convert x and y to tensors
batch_size = 4
shuffle = True
num_workers = 4
training_sample_generator = DataLoader(dataset, batch_size=batch_size, shuffle=shuffle, num_workers=num_workers)

for epoch in range(num_epochs):
    print('Epoch = %s' % epoch)
    for batch_i, samples in enumerate(training_sample_generator):
        features = samples['feature'].float() # Convert features to float
        labels = samples['label'].float() # Convert labels to float

        predictions = model(features)
        error = cost(predictions, labels)
        print('\tBatch = %s, Error = %s' % (batch_i, error.item()))

        optimiser.zero_grad()
        error.backward()
        optimiser.step()

    Epoch = 0
        Batch = 0, Error = 7.702254295349121
        Batch = 1, Error = 13.808805465698242
        Batch = 2, Error = 4.263489246368408
    Epoch = 1
        Batch = 0, Error = 2.072721242904663
        Batch = 1, Error = 1.0039169788360596
        Batch = 2, Error = 0.7218696475028992
    Epoch = 2
        Batch = 0, Error = 1.287321925163269
        Batch = 1, Error = 0.05091605335474014
        Batch = 2, Error = 0.5484017729759216
    Epoch = 3
        Batch = 0, Error = 0.5535609722137451
        Batch = 1, Error = 0.19042828679084778
        Batch = 2, Error = 1.1805700063705444
    Epoch = 4
        Batch = 0, Error = 0.018144484609365463
        Batch = 1, Error = 0.8546347618103027
        Batch = 2, Error = 0.7338478565216064
    Epoch = 5
        Batch = 0, Error = 0.8474566340446472
        Batch = 1, Error = 0.2883010804653168
        Batch = 2, Error = 0.11646117269992828
    Epoch = 6
        Batch = 0, Error = 0.22584526240825653
        Batch = 1, Error = 0.8822470903396606
        Batch = 2, Error = 0.058477893471717834
    Epoch = 7
        Batch = 0, Error = 0.14788007736206055
        Batch = 1, Error = 0.8708349466323853
        Batch = 2, Error = 0.0920991450548172
    Epoch = 8
        Batch = 0, Error = 0.32978326082229614
        Batch = 1, Error = 0.45812496542930603
        Batch = 2, Error = 0.424987256526947
    Epoch = 9
        Batch = 0, Error = 0.7535368800163269
        Batch = 1, Error = 0.19133420288562775
        Batch = 2, Error = 0.0329524502158165

x_for_plotting = np.linspace(0, 2*np.pi, 1000)
design_matrix = torch.tensor(np.vstack([np.ones(x_for_plotting.shape), x_for_plotting])).T, dtype=torch.float32)
print('Design matrix shape:', design_matrix.shape)

y_for_plotting = model.forward(design_matrix)
print('y_for_plotting shape:', y_for_plotting.shape)

    Design matrix shape: torch.Size([1000, 2])
    y_for_plotting shape: torch.Size([1000, 1])

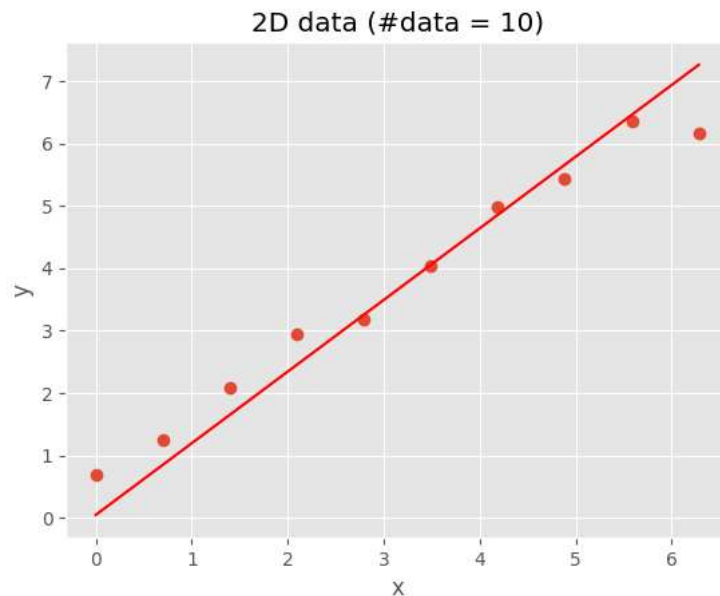
plt.figure()
plt.plot(x, y, 'o')

```

```

plt.plot(x_for_plotting, y_for_plotting.data.numpy(), 'r-')
plt.xlabel('x')
plt.ylabel('y')
plt.title('2D data (#data = %d)' % N)
plt.show()

```



```

def user_cf(M, metric='cosine'):
    pred = np.copy (M)
    n_users, n_items= M. shape
    avg_ratings = np.nanmean (M, axis=1)
    sim_users = sim_matrix(M, 'user', metric)
    for i in range(n_users):
        for j in range(n_items):
            if np.isnan (M[i, j]):
                pred[i, j] = avg_ratings[i] + np.nansum(sim_users[i] * (M[:,j] - avg_ratings)) / sum(sim_users[i])
    return pred

```

```

def item_cf(M, metric='cosine'):
    pred = np.copy (M)
    n_users, n_items = M. shape
    avg_ratings = np.nanmean (M, axis=0)
    sim_items = sim_matrix (M, 'item', metric)
    for i in range(n_users):
        for j in range(n_items):
            if np.isnan (M[i, j]):
                pred[i, j] = avg_ratings[j] + np.nansum(sim_items[j] * (M[i,:] - avg_ratings)) / sum(sim_items[j])
    return pred

```

# Exercise

#1

```

from sklearn.datasets import load_iris
iris = load_iris()

```

```

type(iris)

```

```

sklearn.utils._bunch.Bunch

```

```

iris.data

```

```

array([[5.1, 3.5, 1.4, 0.2],
       [4.9, 3. , 1.4, 0.2],
       [4.7, 3.2, 1.3, 0.2],
       [4.6, 3.1, 1.5, 0.2],
       [5. , 3.6, 1.4, 0.2],
       [5.4, 3.9, 1.7, 0.4],
       [4.6, 3.4, 1.4, 0.3],

```

```
[5. , 3.4, 1.5, 0.2],
[4.4, 2.9, 1.4, 0.2],
[4.9, 3.1, 1.5, 0.1],
[5.4, 3.7, 1.5, 0.2],
[4.8, 3.4, 1.6, 0.2],
[4.8, 3. , 1.4, 0.1],
[4.3, 3. , 1.1, 0.1],
[5.8, 4. , 1.2, 0.2],
[5.7, 4.4, 1.5, 0.4],
[5.4, 3.9, 1.3, 0.4],
[5.1, 3.5, 1.4, 0.3],
[5.7, 3.8, 1.7, 0.3],
[5.1, 3.8, 1.5, 0.3],
[5.4, 3.4, 1.7, 0.2],
[5.1, 3.7, 1.5, 0.4],
[4.6, 3.6, 1. , 0.2],
[5.1, 3.3, 1.7, 0.5],
[4.8, 3.4, 1.9, 0.2],
[5. , 3. , 1.6, 0.2],
[5. , 3.4, 1.6, 0.4],
[5.2, 3.5, 1.5, 0.2],
[5.2, 3.4, 1.4, 0.2],
[4.7, 3.2, 1.6, 0.2],
[4.8, 3.1, 1.6, 0.2],
[5.4, 3.4, 1.5, 0.4],
[5.2, 4.1, 1.5, 0.1],
[5.5, 4.2, 1.4, 0.2],
[4.9, 3.1, 1.5, 0.2],
[5. , 3.2, 1.2, 0.2],
[5.5, 3.5, 1.3, 0.2],
[4.9, 3.6, 1.4, 0.1],
[4.4, 3. , 1.3, 0.2],
[5.1, 3.4, 1.5, 0.2],
[5. , 3.5, 1.3, 0.3],
[4.5, 2.3, 1.3, 0.3],
[4.4, 3.2, 1.3, 0.2],
[5. , 3.5, 1.6, 0.6],
[5.1, 3.8, 1.9, 0.4],
[4.8, 3. , 1.4, 0.3],
[5.1, 3.8, 1.6, 0.2],
[4.6, 3.2, 1.4, 0.2],
[5.3, 3.7, 1.5, 0.2],
[5. , 3.3, 1.4, 0.2],
[7. , 3.2, 4.7, 1.4],
[6.4, 3.2, 4.5, 1.5],
[6.9, 3.1, 4.9, 1.5],
[5.5, 2.3, 4. , 1.3],
[6.5, 2.8, 4.6, 1.5],
[5.7, 2.8, 4.5, 1.3],
[6.3, 3.3, 4.7, 1.6],
[4.9, 2.4, 3.3, 1. ],
```

```
iris.feature_names
```

```
['sepal length (cm)',
'sepal width (cm)',
'petal length (cm)',
'petal width (cm)']
```

```
iris.target
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])
```

Double-click (or enter) to edit

```
iris.target_names
```

```
array(['setosa', 'versicolor', 'virginica'], dtype='<U10')
```

```
type(iris.data)
```

```
numpy.ndarray
```

```
iris.data.shape
```

```

type(iris.target)

        numpy.ndarray

iris.data.shape

        (150, 4)

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

# Load the Iris dataset
iris = load_iris()

# Assign features to X and labels to y
X = iris.data
y = iris.target

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4)

X_train.shape

        (120, 4)

X_test.shape

        (30, 4)

y_train.shape

        (120,)

y_test.shape

        (30,)

from sklearn.linear_model import LinearRegression
from sklearn import metrics

# Assuming you have your training and testing data stored in X_train, X_test, y_train, and y_test

regressor = LinearRegression()
regressor.fit(X_train, y_train)
y_pred = regressor.predict(X_test)

# Evaluate the regression model
mse = metrics.mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2_score = metrics.r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R^2 Score:", r2_score)

        Mean Squared Error: 10.546912775288654
        Root Mean Squared Error: 3.24760108007259
        R^2 Score: 0.8380468731394584

from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
k_range = range(1,20)
scores = {}
scores_list = []
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train,y_train)
    y_pred=knn.predict(X_test)
    scores[k] = metrics.accuracy_score(y_test,y_pred)
    scores_list.append(metrics.accuracy_score(y_test,y_pred))

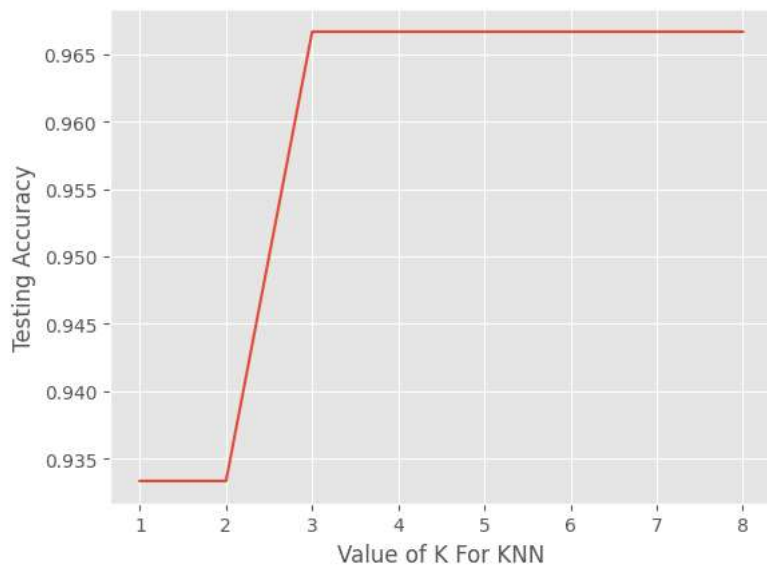
```



```
%matplotlib inline
import matplotlib.pyplot as plt
```

```
plt.plot(k_range,scores_list)
plt.xlabel('Value of K For KNN')
plt.ylabel('Testing Accuracy')
```

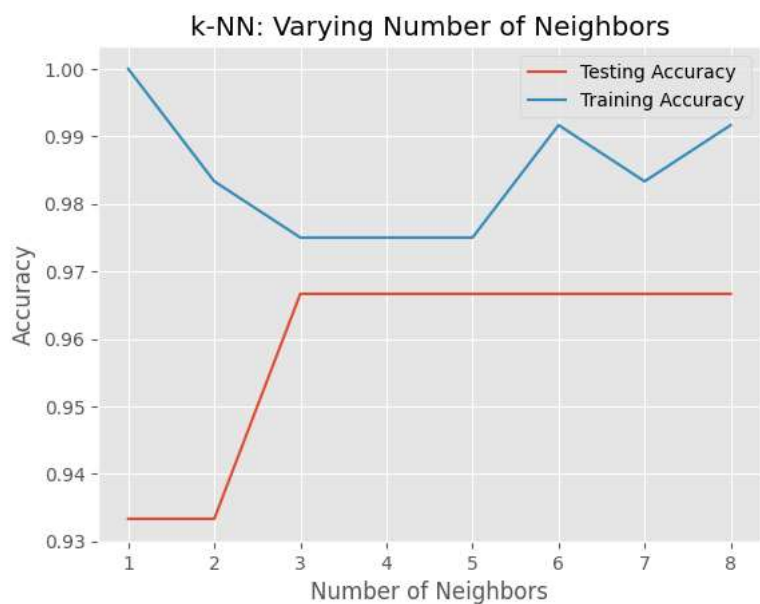
```
Text(0, 0.5, 'Testing Accuracy')
```



```
neighbors = np.arange(1,9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))
```

```
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    train_accuracy[i] = knn.score(X_train, y_train)
    test_accuracy[i] = knn.score(X_test, y_test)
```

```
plt.title('k-NN: Varying Number of Neighbors')
plt.plot(neighbors, test_accuracy, label = 'Testing Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training Accuracy')
plt.legend()
plt.xlabel('Number of Neighbors')
plt.ylabel('Accuracy')
plt.show()
```



#2

#a

#Find list of used genres which is used to category the movies.

```
path1 = 'movies.csv'
path2 = 'ratings.csv'
path3 = 'users.csv'
```

```
movies = pd.read_csv(path1, encoding="ISO-8859-1")
ratings = pd.read_csv(path2)
users = pd.read_csv(path3)
```

```
from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.metrics.pairwise import cosine_similarity
```

```
genres_col = movies['genres']
```

```
genres_list = [genre.split('|') if type(genre) == str else [] for genre in genres_col]
```

```
unique_genres = set([genre for sublist in genres_list for genre in sublist])
```

```
genre_list = list(unique_genres)
```

```
genre_list
```

```
['Documentary',
 'Crime',
 'Musical',
 'Horror',
 'Sci-Fi',
 'Drama',
 'Animation',
 'Romance',
 'Comedy',
 'Adventure',
 'Fantasy',
 'Action',
 "Children's",
 'Thriller',
 'War']
```

#Vectorize the relationship between movies and genres and put them into Ij.

```
unique_genres = sorted(list(set([genre for sublist in genres_list for genre in sublist])))
```

```
Ij = np.zeros((len(movies), len(unique_genres)), dtype=int)
for i, genres in enumerate(genres_list):
    for genre in genres:
        j = unique_genres.index(genre)
        Ij[i, j] = 1
```

```
Ij[:4]
```

```
array([[0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
       [0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0],
       [0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0],
       [0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0]])
```

#Vectorize the relationship between users and genres and put them into Uj (if user rate for a movie, he/she has the related history with the movie)

```
merged_df = pd.merge(ratings, movies, on='movie_id')
unique_user_ids = sorted(list(set(merged_df['user_id'])))
```

```
Uj = np.zeros((len(unique_user_ids), len(unique_genres)), dtype=int)
for i, user_id in enumerate(unique_user_ids):
    user_ratings = merged_df.loc[merged_df['user_id'] == user_id]
    user_genre_list = [genre.split('|') if type(genre) == str else [] for genre in user_ratings['genres']]
    user_genre_set = set([genre for sublist in user_genre_list for genre in sublist])
```

```
for j, genre in enumerate(unique_genres):  
    if genre in user_genre_set:  
        Uj[i, j] = 1
```

```
Uj[:4]
```

```
array([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0],  
       [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1],  
       [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],  
       [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0]])
```